

Multilinear Regression

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

```
In [10]:  import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [20]:  columns = ['Homework', 'Midterm', 'Final']
data = pd.DataFrame({
    "Homework": [95, 70, 80, 100, 70],
    "Midterm": [90, 60, 80, 80, 85],
    "Final": [93, 66, 85, 60, 90]

}, index=['Alice', 'Bob', 'Clare', 'David', 'Eve'])
data.head()
```

Out[20]:

	Homework	Midterm	Final
Alice	95	90	93
Bob	70	60	66
Clare	80	80	85
David	100	80	60
Eve	70	85	90

```
In [21]:  print(np.corrcoef(data['Homework'], data['Final'])[0,1])

-0.1771490677357476
```

```
In [22]:  print(np.corrcoef(data['Midterm'], data['Final'])[0,1])

0.6700743886411277
```

```
In [23]:  theta1 = 0.2
theta2 = 0.8
theta0 = 93 - 0.2*95 - 0.8*90
print(theta0)

2.0
```

```
In [12]:  #MSE
x1 = data.loc['Alice', ['Homework', 'Midterm']].values
print(x1)

[95 90]
```

```
In [24]: y1 = data.loc['Alice',['Final']].values
print(y1)
```

[93]

```
In [25]: theta = np.array([2.0,0.2,0.8])
print(theta)
```

[2. 0.2 0.8]

```
In [26]: prediction = 2.0 + 0.2*x1[0] + 0.8*x1[1]
print(prediction)
```

93.0

```
In [27]: squared_error = (prediction - y1)**2
print(squared_error)
```

[0.]

```
In [28]: def get_squared_error(data,name,theta):
x = data.loc[name,['Homework','Midterm']].values
y = data.loc[name,['Final']].values
prediction = theta[0] + theta[1]*x[0] + theta[2]*x[1]
squared_error = (prediction-y)**2
return squared_error
get_squared_error(data,"Bob",theta)
```

Out[28]: array([4.])

```
In [29]: all_errors = [get_squared_error(data,name,theta)for name in data.index]
print(all_errors)
```

[array([0.]), array([4.]), array([9.]), array([676.]), array([36.])]

```
In [30]: mse = np.mean(all_errors)
print('MSE: ',mse)
print("RMSE: ",np.sqrt(mse))
```

MSE: 145.0

RMSE: 12.041594578792296

Multilinear Regression: Normal Equation

$$\hat{\theta} = (\mathbf{X}^T \cdot \mathbf{X})^{-1} \cdot \mathbf{X}^T \cdot \mathbf{y}.$$

```
In [34]: m,n = data.shape
X = np.hstack([np.ones([m,1]),data[['Homework','Midterm']].values])
print(X)

[[ 1.  95.  90.]
 [ 1.  70.  60.]
 [ 1.  80.  80.]
 [ 1. 100.  80.]
 [ 1.  70.  85.]]
```

```
In [32]: y = data[['Final']].values
print(y)

[[93]
 [66]
 [85]
 [60]
 [90]]
```

```
In [35]: num = X.T.dot(x)
theta_opt = np.linalg.inv(num).dot(X.T.dot(y))
print(theta_opt)

[[35.          ]
 [-0.71627907]
 [ 1.30697674]]
```

```
In [38]: all_errors = [get_squared_error(data,name,theta_opt) for name in data.index]
mse = np.mean(all_errors)
rmse = np.sqrt(mse)
print("MSE: ",mse)
print("RMSE: ",rmse)
```

```
MSE:  36.82790697674422
RMSE:  6.068600083770903
```

Multilinear Regression: Gradient Descent

$$\hat{\theta} \leftarrow \hat{\theta} - r \cdot \frac{\partial J(\hat{\theta})}{\partial \theta}.$$

- The partial derivative of the cost function is given by

$$\frac{\partial J(\hat{\theta})}{\partial \theta} = \frac{2}{m} \cdot \mathbf{X}^T \cdot (\mathbf{X} \cdot \theta - \mathbf{y}).$$

- Verify the formula of partial derivative assuming there is one input feature.**
- End iteration if certain stop criteria is reached, such as:
 - Value of $\hat{\theta}$ becomes stable.
 - Certain iteration amount is reached.

```
In [42]: ▶ m,n=data.shape  
theta_hat = np.random.rand(n,1)  
print(theta_hat)
```

```
[[0.76272655]  
 [0.59038968]  
 [0.14812814]]
```

```
In [47]: ▶ X2 = np.hstack([np.ones([m,1]),data[['Homework','Midterm']].values/100])  
y2 = y  
print(X2)
```

```
[[1.    0.95 0.9 ]  
 [1.    0.7  0.6 ]  
 [1.    0.8  0.8 ]  
 [1.    1.   0.8 ]  
 [1.    0.7  0.85]]
```

```
In [48]: ▶ num_iter = 6000  
r = 0.05  
MSEs = []  
for iter in range(num_iter):  
    gradient = (X2.T).dot(X2.dot(theta_hat)-y2)*2/m  
    theta_hat -= r*gradient  
    MSE = 1/m*(X2.dot(theta_hat)-y2).T.dot(X2.dot(theta_hat)-y2)  
    MSEs.append(MSE[0,0])  
  
print(theta_hat)
```


```
[[ 36.31043372]  
 [-69.94711093]  
 [127.28166283]]
```

```
In [49]: ▶ MSE
```

```
Out[49]: array([[36.92103962]])
```

```
In [50]: ▶ plt.plot(range(num_iter),MSEs)
plt.ylim(0,200)
```

Out[50]: (0, 200)



```
In [51]: ▶ from sklearn.linear_model import LinearRegression
model_lr = LinearRegression()
model_lr.fit(data[['Homework','Midterm']],data['Final'])
```

Out[51]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

```
In [52]: ▶ model_lr.intercept_,model_lr.coef_
```

Out[52]: (35.000000000000004, array([-0.71627907, 1.30697674]))